

Question 1

a. Given X , take $y = 1$ and W_2 such that $W_2^T X = 0$, for any $0 \leq t \leq 1$:

$$(1) \quad E(tW_1 + (1-t)W_2) = (1 - \hat{y}(tW_1))$$

$$(2) \quad tE(W_1) + (1-t)E(W_2) = tE(W_1) + (1-t)\left(1 - \frac{1}{1 + \exp(-0)}\right)^2 = tE(W_1) + (1-t)\left(\frac{1}{2}\right)^2$$

now let $W_1^T X = -N$ where N is a very large number, we have

$$(1) \approx 1$$

and

$$(2) \approx t + \frac{1-t}{4}$$

Take $t = 0.5$ gives

$$(1) > (2)$$

This example shows that $E(W)$ is not necessarily a convex function of W

b. Take $y = 1$, we have

$$(3) \quad E(W) = \log(1 + \exp(-W^T X)) = \log\left(1 + \exp\left(-\sum_{i=1}^N w_i x_i\right)\right)$$

To show (3) a convex function of W , we can simply show the corresponding Hessian matrix is positive semi-definite:

Denote $k = 1 + \exp(-W^T X)$

$$(4) \quad \frac{\partial E(W)}{\partial w_i} = \frac{-x_i \exp(-W^T X)}{k}$$

$$(5) \quad H_{ij} = \frac{\partial^2 E(W)}{\partial w_i \partial w_j} = \exp(-W^T X) \frac{x_i x_j}{k^2}$$

which means the Hessian matrix can be represented as

$$H = \exp(-W^T X) A^T A$$

where

$$A = (w_1 \ w_2 \ \dots \ w_N)$$

thus H is positive semi-definite, we proved what we need.

c. This could be directly seen when we note that no matter $y = 1$ or $y = 0$, we have

$$E(W) = \log(1 + \exp(-W^T X))$$

We can define the loss function as

$$\text{loss}(W) = \sum_{i=1}^N \log(1 + \exp(-W^T X_i))$$

it's a logistics regression problem.

Question 2

a. it's trivial to see that

$$k = \sum_{j=1}^m w_j k_j$$

is symmetric

when k_j is symmetric for all $1 \leq j \leq m$.

Given any vector x

$$(6) \quad x^T k x = \sum_{j=1}^m w_j x^T k_j x$$

as for any j , $w_j \geq 0$, k_j is positive semi-definite, thus

$$w_j x^T k_j x \geq 0$$

we have

$$(6) \geq 0$$

Which implies k is positive semi-definite, thus is a valid kernel.

b. It's trivial to see that K is symmetric let $M^{-1} = \text{diag}(e^{x_1^2}, e^{x_2^2}, \dots, e^{x_n^2})$ then we have

$$(M^{-1})^T K M^{-1} = H = A^T A$$

where $H_{ij} = e^{2x_i x_j}$ and $A = (e^{\sqrt{2}x_1} \ e^{\sqrt{2}x_2} \ \dots \ e^{\sqrt{2}x_n})$ which implies that $K = M^T A^T A M$, thus K is positive semi-definite. So K is a valid kernel.

Question 3

a. In this question, the "best" attribute is decided by its information gain. We have following definition for information gain.

The entropy of a r.v. X with distribution $(p(x_1), p(x_2), \dots, p(x_n))$

$$H(X) = \sum_{i=1}^n -p(x_i) \log_2 p(x_i)$$

The conditional entropy

$$H(X|Y) = \sum_{j=1}^m p(y_j) H(X|y_j)$$

The information gain due to Y

$$IG(X|Y) = H(X) - H(X|Y)$$

Here is the output when running *python dstumpIG.py Mushroom.csv*

```
Fold 0  trainingData errorRate: 0.0164113785558  testData errorRate: 0.0
Fold 1  trainingData errorRate: 0.0164113785558  testData errorRate: 0.0
Fold 2  trainingData errorRate: 0.0164113785558  testData errorRate: 0.0
Fold 3  trainingData errorRate: 0.0164113785558  testData errorRate: 0.0
Fold 4  trainingData errorRate: 0.0164113785558  testData errorRate: 0.0
Fold 5  trainingData errorRate: 0.0150437636761  testData errorRate: 0.012315270936
Fold 6  trainingData errorRate: 0.00724835886214  testData errorRate: 0.0825123152709
Fold 7  trainingData errorRate: 0.011624726477  testData errorRate: 0.0431034482759
Fold 8  trainingData errorRate: 0.0161378555799  testData errorRate: 0.00246305418719
Fold 9  trainingData errorRate: 0.015590809628  testData errorRate: 0.00738916256158
Train Mean ErrorRate: 0.0147702407002  Test Mean ErrorRate: 0.0147783251232
Train StdVar ErrorRate: 0.00287329344241  Test Mean ErrorRate: 0.0258737951366
===== the decision tree =====
# 5 Attribute
|
-- 1 -> Label: 1
|
-- 2 -> Label: -1
|
-- 3 -> Label: -1
|
-- 4 -> Label: -1
|
-- 5 -> Label: 1
|
-- 6 -> Label: 1
|
-- 7 -> Label: -1
|
-- 8 -> Label: -1
|
-- 9 -> Label: -1
```

Our **DTree** class will recursively print out the trained tree

b. In this question, we are using Gini index, which is defined as

$$Gini(X) = \sum_{i \neq j} p(i)p(j)$$

Conditional Gini index

$$Gini(X|Y) = \sum_j p(y_j)Gini(X|y_j)$$

Gini gain

$$\text{GiniGain}(X|Y) = \text{Gini}(X) - \text{Gini}(X|Y)$$

Here is the output when running *python dtree2GI.py Mushroom.csv*

```
Fold 0 trainingData errorRate: 0.00656455142232 testData errorRate: 0.0
Fold 1 trainingData errorRate: 0.00656455142232 testData errorRate: 0.0
Fold 2 trainingData errorRate: 0.00656455142232 testData errorRate: 0.0
Fold 3 trainingData errorRate: 0.00656455142232 testData errorRate: 0.0
Fold 4 trainingData errorRate: 0.00656455142232 testData errorRate: 0.0
Fold 5 trainingData errorRate: 0.00588074398249 testData errorRate: 0.0061576354679
Fold 6 trainingData errorRate: 0.00259846827133 testData errorRate: 0.0357142857143
Fold 7 trainingData errorRate: 0.00574398249453 testData errorRate: 0.0073891625615
Fold 8 trainingData errorRate: 0.00629102844639 testData errorRate: 0.0024630541871
Fold 9 trainingData errorRate: 0.00574398249453 testData errorRate: 0.0073891625615
Train Mean ErrorRate: 0.00590809628009 Test Mean ErrorRate: 0.00591133004926
Train StdVar ErrorRate: 0.0011536704983 Test Mean ErrorRate: 0.0103887175906
===== the decision tree =====
# 5 Attribute
|
-- 1 ->
    # 1 Attribute
    |
    -- 2 -> Label: 1
    |
    -- 3 -> Label: 1
    |
    -- 6 -> Label: 1
|
-- 2 ->
    # 1 Attribute
    |
    -- 6 -> Label: -1
|
-- 3 ->
    # 1 Attribute
    |
    -- 3 -> Label: -1
    |
    -- 4 -> Label: -1
    |
    -- 6 -> Label: -1
|
-- 4 ->
    # 1 Attribute
    |
    -- 3 -> Label: -1
```

```

      |
      -- 4 -> Label: -1
      |
      -- 6 -> Label: -1
|
-- 5 ->
    # 1 Attribute
    |
    -- 2 -> Label: 1
    |
    -- 3 -> Label: 1
    |
    -- 6 -> Label: 1
|
-- 6 ->
    # 19 Attribute
    |
    -- 1 -> Label: 1
    |
    -- 2 -> Label: 1
    |
    -- 3 -> Label: 1
    |
    -- 4 -> Label: 1
    |
    -- 5 -> Label: 1
    |
    -- 6 -> Label: -1
    |
    -- 8 -> Label: 1
    |
    -- 9 -> Label: 1
|
-- 7 ->
    # 1 Attribute
    |
    -- 3 -> Label: -1
    |
    -- 6 -> Label: -1
|
-- 8 ->
    # 1 Attribute
    |
    -- 3 -> Label: -1
    |

```

```

-- 4 -> Label: -1
|
-- 6 -> Label: -1
|
-- 9 ->
# 1 Attribute
|
-- 3 -> Label: -1
|
-- 4 -> Label: -1
|
-- 6 -> Label: -1

```

Question 4

a. AdaBoost using the additive model, it assumes that the final classifier could be represented as

$$g(x) = \text{sign}\left[\sum_{t=1}^T \alpha_t G_t(x)\right]$$

it is trying to minimize the following error function

$$E = \sum_{i=1}^N \exp(-y G_t(x_i))$$

suppose the initial distribution on data X is W , then we have the following update scheme, for $t = 1, 2, \dots, T$:

- try to obtain a G_t in the hypothesis space that minimize $\epsilon_t = \Pr_{x \sim W_t}[G_t(x) \neq y]$
- choose $\alpha_t = \frac{1}{2} \left(\frac{1 - \epsilon_t}{\epsilon_t} \right)$
- update $W_{t+1} = \frac{W_t(i) \exp(-\alpha_t y_i G_t(x_i))}{Z_t}$

where Z_t is a normalization factor.

Following is the input and output:

```
$ python myAdaBoost.py Mushroom.csv 5
```

```
T = 5
```

```

Fold 0  trainingData errorRate:  0.0164113785558  testData errorRate: 0.0
Fold 1  trainingData errorRate:  0.0164113785558  testData errorRate: 0.0
Fold 2  trainingData errorRate:  0.0164113785558  testData errorRate: 0.0
Fold 3  trainingData errorRate:  0.0164113785558  testData errorRate: 0.0
Fold 4  trainingData errorRate:  0.0164113785558  testData errorRate: 0.0
Fold 5  trainingData errorRate:  0.0150437636761  testData errorRate: 0.012315270936
Fold 6  trainingData errorRate:  0.00793216630197  testData errorRate: 0.0665024630542
Fold 7  trainingData errorRate:  0.011624726477   testData errorRate: 0.0431034482759
Fold 8  trainingData errorRate:  0.0161378555799  testData errorRate: 0.00246305418719
Fold 9  trainingData errorRate:  0.011761487965   testData errorRate: 0.0221674876847
Train Mean ErrorRate: 0.0144556892779  Test Mean ErrorRate: 0.0146551724138

```

Train StdVar ErrorRate: 0.00283069851513 Test Mean ErrorRate: 0.0218674966678

\$ python myAdaBoost.py Mushroom.csv 10

T = 10

Fold	trainingData errorRate:	testData errorRate:
Fold 0	0.0103938730853	0.0
Fold 1	0.0103938730853	0.0
Fold 2	0.00929978118162	0.0
Fold 3	0.0103938730853	0.0
Fold 4	0.0175054704595	0.0
Fold 5	0.0138129102845	0.0135467980296
Fold 6	0.00793216630197	0.0665024630542
Fold 7	0.00998358862144	0.0480295566502
Fold 8	0.0150437636761	0.00246305418719
Fold 9	0.0166849015317	0.00738916256158
Train Mean ErrorRate: 0.0121444201313		Test Mean ErrorRate: 0.0137931034483
Train StdVar ErrorRate: 0.0031661392124		Test Mean ErrorRate: 0.0225191051623

\$ python myAdaBoost.py Mushroom.csv 20

T = 20

Fold	trainingData errorRate:	testData errorRate:
Fold 0	0.0109409190372	0.0
Fold 1	0.0103938730853	0.0
Fold 2	0.0120350109409	0.0
Fold 3	0.0120350109409	0.0
Fold 4	0.0158643326039	0.0172413793103
Fold 5	0.0138129102845	0.0135467980296
Fold 6	0.00615426695842	0.0628078817734
Fold 7	0.00765864332604	0.0295566502463
Fold 8	0.0106673960613	0.00246305418719
Fold 9	0.00683807439825	0.012315270936
Train Mean ErrorRate: 0.0106400437637		Test Mean ErrorRate: 0.0137931034483
Train StdVar ErrorRate: 0.00291324441044		Test Mean ErrorRate: 0.018853227816

\$ python myAdaBoost.py Mushroom.csv 40

T = 40

Fold	trainingData errorRate:	testData errorRate:
Fold 0	0.0125820568928	0.0
Fold 1	0.0103938730853	0.0
Fold 2	0.0103938730853	0.0
Fold 3	0.0103938730853	0.0
Fold 4	0.0158643326039	0.0172413793103
Fold 5	0.0138129102845	0.0135467980296
Fold 6	0.00533369803063	0.0431034482759
Fold 7	0.0082056892779	0.0295566502463
Fold 8	0.00902625820569	0.0024630541871

```
Fold 9  trainingData errorRate: 0.00902625820569  testData errorRate: 0.012315270936
Train Mean ErrorRate: 0.0105032822757  Test Mean ErrorRate: 0.0118226600985
Train StdVar ErrorRate: 0.00283581466702  Test Mean ErrorRate: 0.0140437293499
```

b. For LogitBoost, There is no big change comparing with AdaBoost. As pointed out in [1], the only modification is to let the distribution $W_t(i)$ be proportional to

$$\frac{1}{1 + \exp(y_i f_{t-1}(x_i))}$$

where $f_t = \sum_{j=1}^t \alpha_j G_j$ is the current trained model. The detail of derivation could be found in [2]. But it is a little bit beyond my ability to fully understand it..

c. Algorithm for this question was described in Question 4.(b). Here is the input and output

```
$ python myLogitBoost.py Mushroom.csv 5
```

```
T = 5
```

```
Fold 0  trainingData errorRate: 0.0164113785558  testData errorRate: 0.0
Fold 1  trainingData errorRate: 0.0164113785558  testData errorRate: 0.0
Fold 2  trainingData errorRate: 0.0164113785558  testData errorRate: 0.0
Fold 3  trainingData errorRate: 0.0164113785558  testData errorRate: 0.0
Fold 4  trainingData errorRate: 0.0164113785558  testData errorRate: 0.0
Fold 5  trainingData errorRate: 0.0150437636761  testData errorRate: 0.012315270936
Fold 6  trainingData errorRate: 0.00724835886214  testData errorRate: 0.0825123152709
Fold 7  trainingData errorRate: 0.011624726477  testData errorRate: 0.0431034482759
Fold 8  trainingData errorRate: 0.0161378555799  testData errorRate: 0.00246305418719
Fold 9  trainingData errorRate: 0.015590809628  testData errorRate: 0.00738916256158
Train Mean ErrorRate: 0.0147702407002  Test Mean ErrorRate: 0.0147783251232
Train StdVar ErrorRate: 0.00287329344241  Test Mean ErrorRate: 0.0258737951366
```

```
$ python myLogitBoost.py Mushroom.csv 10
```

```
T = 10
```

```
Fold 0  trainingData errorRate: 0.0164113785558  testData errorRate: 0.0
Fold 1  trainingData errorRate: 0.0164113785558  testData errorRate: 0.0
Fold 2  trainingData errorRate: 0.0164113785558  testData errorRate: 0.0
Fold 3  trainingData errorRate: 0.0164113785558  testData errorRate: 0.0
Fold 4  trainingData errorRate: 0.0164113785558  testData errorRate: 0.0
Fold 5  trainingData errorRate: 0.0150437636761  testData errorRate: 0.012315270936
Fold 6  trainingData errorRate: 0.00724835886214  testData errorRate: 0.0825123152709
Fold 7  trainingData errorRate: 0.011624726477  testData errorRate: 0.0431034482759
Fold 8  trainingData errorRate: 0.0161378555799  testData errorRate: 0.00246305418719
Fold 9  trainingData errorRate: 0.015590809628  testData errorRate: 0.00738916256158
Train Mean ErrorRate: 0.0147702407002  Test Mean ErrorRate: 0.0147783251232
Train StdVar ErrorRate: 0.00287329344241  Test Mean ErrorRate: 0.0258737951366
```



```
$ python myLogitBoost.py Mushroom.csv 20
```

```
T = 20
```

```
Fold 0 trainingData errorRate: 0.0164113785558 testData errorRate: 0.0
Fold 1 trainingData errorRate: 0.0164113785558 testData errorRate: 0.0
Fold 2 trainingData errorRate: 0.0164113785558 testData errorRate: 0.0
Fold 3 trainingData errorRate: 0.0164113785558 testData errorRate: 0.0
Fold 4 trainingData errorRate: 0.0164113785558 testData errorRate: 0.0
Fold 5 trainingData errorRate: 0.0150437636761 testData errorRate: 0.012315270936
Fold 6 trainingData errorRate: 0.00724835886214 testData errorRate: 0.0825123152709
Fold 7 trainingData errorRate: 0.011624726477 testData errorRate: 0.0431034482759
Fold 8 trainingData errorRate: 0.0161378555799 testData errorRate: 0.00246305418719
Fold 9 trainingData errorRate: 0.015590809628 testData errorRate: 0.00738916256158
Train Mean ErrorRate: 0.0147702407002 Test Mean ErrorRate: 0.0147783251232
Train StdVar ErrorRate: 0.00287329344241 Test Mean ErrorRate: 0.0258737951366
```

```
$ python myLogitBoost.py Mushroom.csv 40
```

```
T = 40
```

```
Fold 0 trainingData errorRate: 0.0164113785558 testData errorRate: 0.0
Fold 1 trainingData errorRate: 0.0164113785558 testData errorRate: 0.0
Fold 2 trainingData errorRate: 0.0164113785558 testData errorRate: 0.0
Fold 3 trainingData errorRate: 0.0164113785558 testData errorRate: 0.0
Fold 4 trainingData errorRate: 0.0164113785558 testData errorRate: 0.0
Fold 5 trainingData errorRate: 0.0150437636761 testData errorRate: 0.012315270936
Fold 6 trainingData errorRate: 0.00724835886214 testData errorRate: 0.0825123152709
Fold 7 trainingData errorRate: 0.011624726477 testData errorRate: 0.0431034482759
Fold 8 trainingData errorRate: 0.0161378555799 testData errorRate: 0.00246305418719
Fold 9 trainingData errorRate: 0.015590809628 testData errorRate: 0.00738916256158
Train Mean ErrorRate: 0.0147702407002 Test Mean ErrorRate: 0.0147783251232
Train StdVar ErrorRate: 0.00287329344241 Test Mean ErrorRate: 0.0258737951366
```

References

- [1] Schapire, Robert E. *"The boosting approach to machine learning: An overview."* LECTURE NOTES IN STATISTICS-NEW YORK-SPRINGER VERLAG- (2003): 149-172.
- [2] Lafferty, Guy Lebanon John. *"Boosting and maximum likelihood for exponential models."* Advances in neural information processing systems 1 (2002): 447.